

CS 61C Final Review

Logistics

Scope

- all lectures up to CALL (Mon 9/27)
- disc 1-5
- labs 1-4
- hw 1-4
- 1 and 2a

Topical Review with Questions

- Number Representation (Avi)
- C (Ryan)
 - Basics
 - Pointers
 - Arrays
- Memory Management (Ryan)
- Floating Point (Ryan)
- RISC-V Basics (Avi)
- RISC-V Coding (Avi)
- Instruction Formats (Avi)
- CALL (Avi)

Number Representation

- common questions are just converting between bases
- Fall 2018 Q1b
 - binary to decimal: 0b01101101 \rightarrow 109

Example

- 26_{10} to binary
 - $16+8+2$ so 0b11010
- 26_{10} to hex
 - pad binary with leading 0s to fill nibble
 - 0b0001+1010 = 0x1A

Quick Tricks

- how to represent power of two in binary
 - only 1 bit is 1, the rest are 0
- representing one less than a power of 2
 - all ones up to the next power of 2
- Two's Complement
 - like normal except you start with all 1s and use 0s to negate

Two's Complement: Negative Nums

- treat most sig bit (leftmost_ as adding a negative power of two

Example: interpreting this number

	<u>1</u>	<u>0</u>	<u>1</u>	<u>0</u>	<u>1</u>	<u>1</u>	<u>0</u>	<u>1</u>
8-bit unsigned:	2^7		2^5		2^3	2^2		$2^0 = 128 + 32 + 8 + 4 + 1 = 173$
8-bit 2's comp:	-2^7		2^5		2^3	2^2		$2^0 = -128 + 32 + 8 + 4 + 1 = -83$

Converting Binary to Hex

0b10100

- hex is implicitly unsigned
- so converting a negative number to hexadecimal is hard
- -64, first convert to binary, then convert binary to hex
- just (mod 16) and take the remainders!
 - fill in from least significant bit to most
 -

Bias Notation

- fixes the problem of converting negative nums to hex
- in bias notation we shift the unsigned interpretation of the number
- 8-bit with a bias of -253
 - smallest representatable num = 0b0000_0000 is $0 - 252 = -252$
 - largest num = 0b1111_1111 = $255 - 252 = 3$
- we don't have to put the bias in the middle!
- bias of 0 is just a regular unsigned bit

Fall 2019 Q1

$$(37)_{10} = (25)_{16}$$

Step by step solution

Step 1: Divide $(37)_{10}$ successively by 16 until the quotient is 0:

$$37/16 = 2, \text{ remainder is } 5$$

$$2/16 = 0, \text{ remainder is } 2$$

Step 2: Read from the bottom (MSB) to top (LSB) as 25. So, 25 is the hexadecimal equivalent of decimal number 37 (Answer).

Decimal:	0	1	2	3	4	5	6	7
Hexa	0	1	2	3	4	5	6	7
Decimal:	8	9	10	11	12	13	14	15
Hexa	8	9	A	B	C	D	E	F

Negate the following nibble binary/hex numbers, or write N/A if not possible. Remember to write your answer in the appropriate base.

Representation	Binary/Hex	Negation
Unsigned	0b0101	N/A (no negatives in unsigned)
Bias = -7	0b0100	0b1010
Bias = -7	0xF	N/A (out of range)
Two's Complement	0b1100	0b0100
Two's Complement	0xA	0x6

- 0xF = 15 so out of range!
- 0b0100 = 8, bias of -7 = 0b1010

CHEAT SHEET STUFF

Cheat Sheet Stuff

- N bits represents 2^N “things”
 - N-bit unsigned: $[0, 2^N - 1]$
 - Smallest: all 0s, base 10 value 0
 - Largest: all 1s, base 10 value $2^N - 1$
 - N-bit two's complement: $[-2^{N-1}, 2^{N-1} - 1]$
 - Smallest: 1 in sign bit followed by all 0s, base 10 value -2^{N-1}
 - Largest: 0 in sign bit followed by all 1s, base 10 value $2^{N-1} - 1$
 - base 10 value of -1 is represented by all 1s
 - To negate: flip all bits and add 1
 - ^ think about why for all of these!
 - N-bit bias notation with bias b : $[b, 2^N - 1 + b]$
 - Smallest: all 0s, base 10 value b
 - Largest: all 1s, base 10 value $2^N - 1 + b$
- Glance over definitions of “sign and magnitude” from lecture slides (less important than the three above notations, but will come up occasionally)

Bitwise Operations

- $\&$ = AND = turns bits ON
- $|$ = OR = turns bits OFF
- \wedge = XOR = flips bits

C Types

- C is weakly typed (casting is hard)

Types

- char (8 bits, 1 byte)
- int (32 bits, 4 bytes)
 - `int32_t`, `uint8_t`, etc. have explicit sizes given by `<stdint.h>`
- `int*` (always either 32 bit or 64 bit based on operating system)

False Values

- `'\0'` = null terminator of char
- 0
- NULL pointer

Pointer Arithmetic

If you do arithmetic with pointers, remember that you're incrementing in `sizeof(pointer_type)`, not bytes!

Example: assume `sizeof(int) == 4`

```
int x = 5;
int *y = &x;           // y = 0x4000
y = y + 2;             // y = 0x4008
char* z = "bears!";    // z = 0x2000
z = z + 2;             // z = 0x2002
```

Bit Manipulation

- review lab 2 exercise 3!!

Example (Fa19 Q3)

```
1 void ConvertTo2sArray(int32_t *A) {
2     while (*A) {           // checks if A != 0x00000000 (terminator)
3         if (*A < 0) {      // checks if value in A is negative
4             ConvertTo2s(A); // calls other function on address
5         }
6         A = A + 1;         // increments A's address by 1
7     }
8 }
9
10 void ConvertTo2s(int32_t *B) {
11     *B = (*B & 0x7FFF_FFFF) + 1; // converts to 2s complement
12 }
13
```

Structs

- kinda like classes
- also just some sequence of bits

sizeof

- returns num of bits being occupied
- `sizeof` an array and `sizeof` a pointer are diff based on the size of the array and the platform you are on

```

1 char x[5] = {0, 0, 0, 0, 0};
2 char *y = calloc(5, sizeof(char));
3 //size of
4 --> sizeof(x) == 5 * sizeof(char) == 5
5 --> sizeof(y) == sizeof(char *) == 4
6 // (or 8 depending on your platform)

```

Memory

Endianness

- bucket of buckets!
- if we have data 0x01020304:

Address	0x100	0x101	0x102	0x103
Little Endian	0x04	0x03	0x02	0x01
Big Endian	0x01	0x02	0x03	0x04

Types

- **Stack** = function local variables, strings allocated as arrays
- **Heap** = dynamically allocated memory (with malloc, calloc, realloc)
- **Static** = global variables, statically allocated strings
- **Code** = machine instructions

Example

```

1 int a = 5;
2 int main(){
3     int b = 0;
4     char* s1 = "cs61c";
5     char s2[] = "cs61c";
6     char *c = malloc(sizeof(char) * 100);
7     return 0;
8 }

```

- s1 = stack
- s2 = stack
- s1[0] = static
- s2[0] = stack
- c[0] = heap
- a = static

Practice

Floating Point

Practice Problems

- What's the largest representable number?
- What's the smallest representable number?
- What's the smallest unrepresentable number?
- What's the largest unrepresentable numbers?
- What's the smallest distance between two numbers? What about denorms?
- How do you represent a particular number in floating point?
 - It basically reduces to these last two (see discussion 3 for practice with other types of problems)

Conversion

Sign	Exponent	Mantissa/Significand/Fraction
1	8	23

- **Normalized** = $(-1)^{Sign} * 2^{Exp+Bias} * 1.significand_2$
- **Denormalized** = $(-1)^{Sign} * 2^{Exp+Bias+1} * 0.significand_2$

Exponent	Significand	Meaning
0	0	0
0	any	denorm
255	0	+ or - ∞
255	nonzero	NaN
1-254	any	norm

Example Fa18 Q1

- **minifloat** = 8-bits with 1sign, 5exp, 2mantissa, bias -15
 - SEEEMMM

- how many NaNs do we have?
 - we have **6** because we need max exponent and nonzero mantissa bits
 - $2 * (2^2 - 1) = 6$
 - **bit rep in hex of next minifloat bigger than 0x3F**
 - 0x3F = 0b0011_0111
 - which is SEEEEEEMMM
 - increasing mantissa would cause overflow, so we need to increment the exponent once and wrap mantissa back to 0!
 - we get 0b0100_0000 = 0x40 (or just $0x3F + 1$)
 - **bit rep of encoding -2.5**
 - convert to binary: $-2.5 = -10.1_2$
 - sign = 1
 - $\text{exp} + \text{bias} = \text{exp} - 15 = 1$, so $\text{exp} = 16 = 10000$
 - $.5 = 10$
 - 1. represent num in binary
 1. $-2.5 = -10.1_2$
 2. normalize so that one number is before the binary point
 1. $-10.1 = -1.01 * 2^1 \rightarrow \text{MM} = 01$
 2. Ex: $1101 \rightarrow 1.101 * 2^3$
 3. calculate exp bits
 1. $1 = \text{EEEEEE} - 15 = 10000$ (we need 16)
 4. sign bit
 1. 1 = negative
 2. 0 = positive
 5. Put it together!
 1. $1100_0001 = 0xC1$
- **What does this return?**

```

1 // NOTE: always round down to 0
2 minifloat should_be_a_billion() {
3     minifloat sum = 0.0;
4     for (unsigned int i = 0; i < 1000000000; i++) {
5         sum = sum + 1.0;
6     }
7     return sum;
8 }

```

- if you have a really big number, the FP can't represent too small or too big of a number
 - you need to find the largest number that the minifloat can represent
- once the $\text{exp} = 3$, you would need to increase the sum by 2.0 to actually change the number
- When does incrementing mantissa increment by 1.0?
 - at $1.11 * 2^2 = 111 = 7$, so we would need the exponent to be 2

- if the exponent gets larger, the represented number is $1.00 * 2^3 = 1000 = 8$
- **shorthand way** = once the exponent is greater than the number of mantissa bits, the number will start increasing by more than just 1.0 every time.
- so 8.0 is returned!

RISC-V

C to RISC-V

- C constructs are translated to assembly
- like while loops!

<pre>int s0 = 0; while (s0 < 4) { do_work(); s0 += 1; }</pre>	<pre>li s0, 0 while: li t0, 4 bge s0, t0, done jal do_work addi s0, s0, 1 j while done:</pre>
--	---

- sw saves registers to memory
- lw loads memory to registers

Jumping

	labels	registers
no linking	j	jr
linking	jal	jar

Calling Convention

- make sure that when you are writing a function, you HAVE TO save your save registers and then restore them!
- temp variables can CHANGE if you jal to something else :')
- you need to save ur return address (ra) to the stack so you know how to get back when you jump to a different function

Example Su19 Q5

- implement strcpy in RISC-V

- `char* strncpy(char* destination, char* source, unsigned int n);`
- takes in two `char*` args and copies up to the first `n` characters
- if it reaches a null terminator, it copies that value into the destination and then stops
- returns a pointer to the destination string

```

1 # iterative python-ish pseudocode:
2 i = 0
3 while (i != n):
4     copy_one_char_from_source_to_dest
5     if source_is_null_terminator:
6         return dest
7     increment source pointer
8 return dest

```

```

1 # iterative RISC-V solution
2 # a0 = dest, a1 = source, a2 = n
3 strncpy:
4     li t0, 0      # initialize counter
5     mv t2, a0     # store a0 in t2 as a temp
6 loop:
7     beq t0, a2, end # check end of while loop
8     lb t1, 0(a1)   # read char from source
9     sb t1, 0(a0)   # copy char to dest
10    beq t1, x0, end # check null terminator
11    addi t0, t0, 1  # increment counter
12    addi a1, a1, 1  # increment source
13    addi a0, a0, 1  # increment dest
14    j loop         # go back to top of loop
15 end:
16    mv a0, t2      # t2 contains the original dest
17    ret

```

- since we are loading and saving ints, we just need to do `lb` and `sb` since ints are just 1 byte
 - `lw` and `sw` depend on the system
- there's also a way to do it recursively!

```

1 # recursive solution
2 # a0 = dest, a1 = source, a2 = n
3 strncpy:
4     # missing! we need to save the return address and s0
5     addi sp, sp, -8
6     sw ra, 0(sp)
7     sw s0, 4(sp)
8
9     beq a2, x0, end # n == 0
10    lb s0, 0(a1)    # read *source
11    sb s0, 0(a0)    # copy to *destination

```

```

12  beq s0, x0, end # check if *source == 0
13  # set up recursive call
14  addi a0, a0, 1
15  addi a1, a1, 1
16  addi a2, a2, -1
17  jal strncpy
18  addi a0, a0, -1 # decrements the address back to the destination
19  end:
20  # missing!
21  lw ra 0(sp) # restores ra and s0
22  lw s0 4(sp)
23  addi sp sp 8
24
25  ret # return (a0 is already dest)
26  # Q: What's missing?
27  # A: Calling convention (need to save s0 and ra)!

```

- CC contract
 - says that ra = contains the address to jump back to
 - and that sp is the same before AND after a function call
 - as long as these two rules are followed, everything works!

Instruction Formats

- SB = (B)ranches = conditional jumps
- U = (U)pper Immediates = long immediates
- UG = (J)umps = unconditional jumps
- R = (R)egister
- I = (I)mmediate + loads from memory
- S = (S)tores to memory



Practice Problem

lw t5, 17(t6)

- lw = 0b0000011 (I-type)
- funct3 = 010
- t5 = x30 = 0b11110
- t6 = x31 = 0b11111
 - 17 (in 12 bits) = 0b0000_0001_0001
- For I types
 - imm[11:0] rs1 funct3 rd opcode
 - 0000_0001_0001 1111 010 11110 0000011
 - 0b0000_0001_0001_1111_1010_1111_0000_0011

- or 0x011FAF03
- **DON'T mess up converting to hex!!!!**

CALL

C = Compile

- C -> risc or other assembly languages
- produced assembly is similar to what you might write by hand

A = Assemble

- RISC-V -> object files
- expands psuedo instructions
- requires 2 passes to resolve labels within files
- produces machine code with relocation tables (identifies labels defined in *other* files) and symbol tables (identifies labels/addresses that *this* file defines)

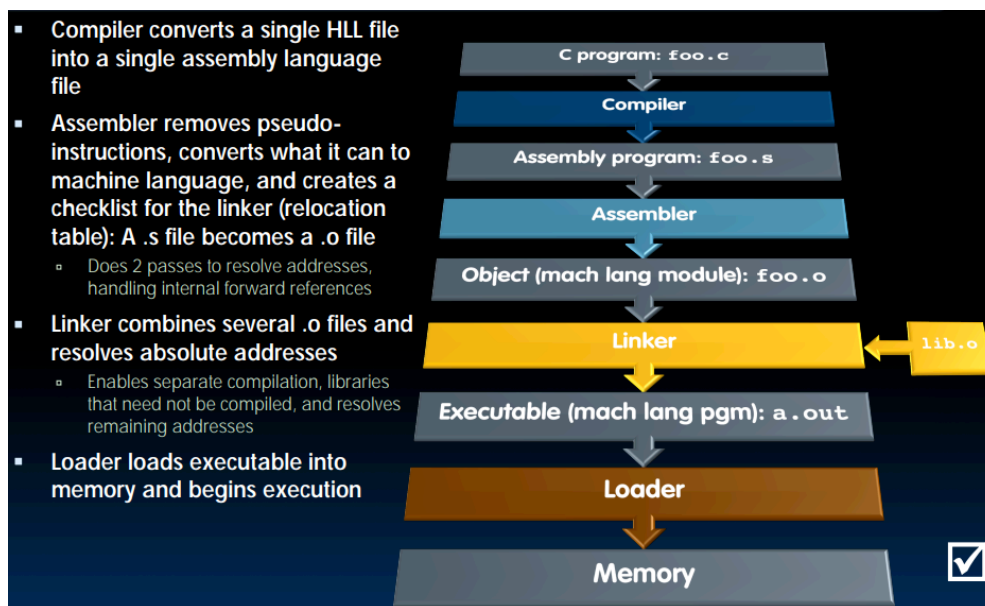
L = Link

- combines object files together into binary executable
- fulfills missing labels/addresses required by examining relocation and symbol tables

L = Load

- part of operating system
- loads machine code into memory for execution

Summary of CALL

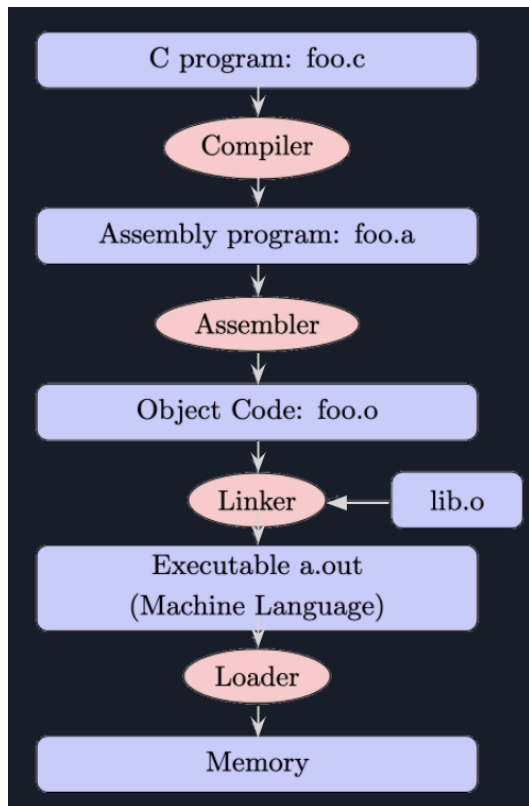


Lecture 13: CALL Notes

Interpretation vs Translation

- **Interpreter** = runs and executes program directly
 - python, java, c++ = easy to program, inefficient to interpret
 - assembly, machine code, bytes = difficult to program, efficient!
 - easier to write interpreter
 - closer to high-level, so better error and debugging messages
 - interpreters are slower and smaller
 - provides instruction set independence (runs on any machine)
 - usually open source
- **Translator** = converts a program from source language to equivalent program in another language and then runs it
 - helps to “hide” the program/intellectual property from its users
 - translates and compiles to a different machine!

Full Picture



- see the big picture by doing `gcc -O2 -S -c foo.c`, the `-S` makes it verbose

Compiler

- input = high-level language code (like foo.c)
- output = assembly language code (like foo.s for RISC-V)
- output may contain pseudo-instructions
 - instructions that assembler understands but not machine (like “copy value from t2 to t1”)
 - `mv t1 t2` → `addi t1 t2 0`
 - assembler added things like `mv` so we don't have to use simple commands like `addi`

Assembler

- input = assembly code (like foo.s for RISC-V)
- output = object code, info tables (true assembly only)
 - e.g., foo.o for RISC-V
- reads and uses directives
- replaces pseudo-instructions with machine language and then creates object file

Assembler Directives

gives directions to assembler, but doesn't produce machine instructions

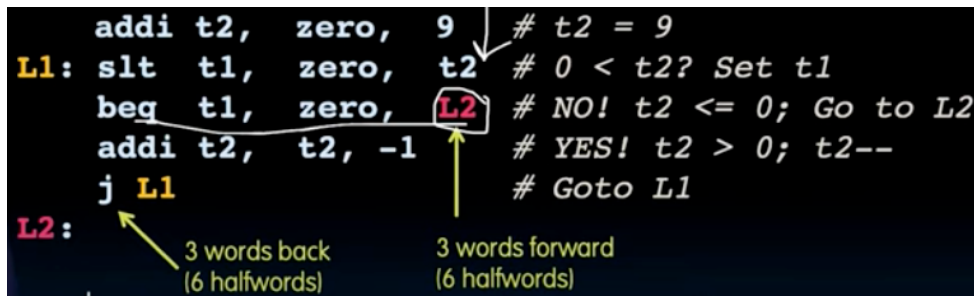
- `.text` = subsequent items put in user text segment (machine code)
- `.data` = subsequent items put in user data segment (source file data in binary)
- `.global sym` = declares `sym` global and can be references from other files
- `.string str` = stores string `str` in memory and null-terminates it
- `.word w1...wn` = store the `n` 32-bit quantities in successive memory words

Pseudo-instruction Replacement

- `neg t0 t1` is the same as `sub t0 zero t1`
- `la t0 str`
 - `lui t0 str[31:12]` then `addi t0 t0 str[11:0]` (STATIC addressing)
 - `auipc t0 str[31:12]` then `addi t0 t0 str[11:0]` (PC-relative addressing)
- don't forget!!!!
 - sign extended immediates
 - machine code can only do operations on 32 bits, so if you have 12 bits, you have to sign-extend it to 32 bits
 - Branch imms count halfwords
 - each line is a “word”, so if we wanted to travel back 3 lines, we would go back 6 halfwords

Producing Machine Code

- simple case
 - arithmetic, logical, shifts, etc.
- branches and jumps are PC-relative (e.g., beq/bne, jal)
- once pseudo-instructions are replaced by real ones, we know by how many instructions to branch/jump over
- “Forward Reference” problem
 - branch instructions sometimes branch to labels that are forward in the problem (like a loop)
 - SOLVED: take two passes over the program
 - first = remember position of labels
 - second = use label positions to generate the code
 - Example of halfwords and remembering label positions:



- for PC-relative jumps (jal) and branches (bne, ben, etc)
 - j offset (pseudo-instructions that expands to `jal zero, offset`)
 - count the number of instruction “half words” between the target and jump to determine the offset (position-independent code aka PIC)
- references to static data
 - `la` gets broken into `lui` and `addi`
 - requires the full 32-bit address of data

Tables

- **Symbol table** = list of “items” in this file that may be used by other files
 - labels = function calling
 - data = anything in the `.data` section, variables that may be accessed across files
- **Relocation table** = list of “items” whose address this file needs
 - basically saying here’s something to fill in the blank, what do we need to remember/store?
 - any absolute label jumped to (using `jal`, `jalr`)
 - internal
 - external (like lib files)
 - EX: the `la` instruction (e.g., for `jalr` base register)
 - any piece of data in static section

- EX: the `la` instruction (e.g., for `lw/sw` base register)

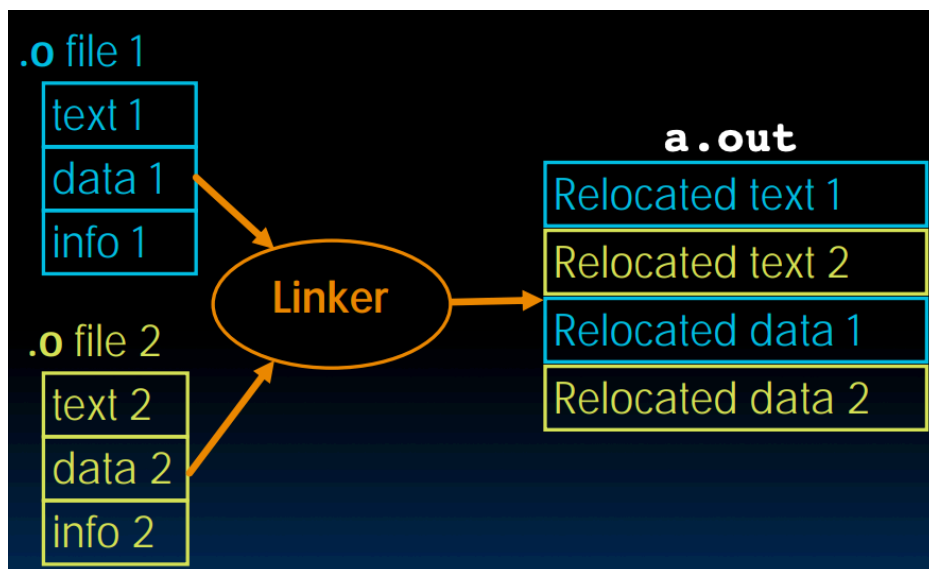
Object File Format

- **object file header** = size and position of the other pieces of the object file
- **text segment** = the machine code
- **data segment** = binary representation of the static data in the source file
- **relocation information** = identifies lines of code that need to be fixed up later
- **symbol table** = list of this file's labels and static data that can be referenced
- **debugging information**

Linker

General Info

- input = object code files, info tables (like `foo.o`, `libc.o`)
- output = executable code (`a.out` for RISC)
- combines several object (`.o`) files into a single executable aka "linking" them together
- compiles files, BUT doesn't have to recompile if you change one file!



Steps

1. take *text* segment from each `.o` file and put them together
2. take *data* segment from each `.o` file, put them together, and concatenate onto the end of text segments
3. resolve references (i.e., handle each entry from reloc table, fill in all absolute addresses)

4 Types of Addresses

- PC-Relative Addressing
 - `ben`, `bne`, `jal`, `auipc/addi`
 - never need to relocate (cuz PIC aka Position-Independent Code)
- Absolute Function Address

- auipc/jalr
 - always relocate
- External Function Reference
 - auipc/jalr
 - always relocate
- Static Data Reference
 - lui/addi
 - always relocate

Relocation Editing

- J-format = jump/jump and link
 - the upper bits need to be modified
- I-, S- = loads and stores to variables in static area
 - all relative to global pointer so all the upper bits need to be modified based on where the static area begins
- conditional branches don't need to be modified since they are PC-relative!

Resolving References

- the linker assumes the first word of the first text segment is at 0x10000
- it knows the length and ordering of each text and data segment
- it calculates the absolute address of each label to be jumped to and each data being referenced
- to resolve!
 - search for reference (data or label) in all user *symbol tables*
 - if not, search the library files
 - once absolute address is determined, fill in the machine code appropriately
- *conflicting symbol type* = occurs when the machine finds two of the same symbol
- output = executable file containing text, data, and header

Static vs Dynamically Linked Libraries

- what we just described is statically linked
 - library is part of the executable ("baked in")
- dynamically-linked libraries (DLL)
 - executable doesn't have to recompile every time
 - puts the responsibility on the loader to get the library when it runs
 - space/time issues
 - storing a program requires less disk space
 - sending a program requires less time
 - 2 programs requires less memory since they share a library

- upgrades
 - replacing one library file upgrades every program that uses the library
 - but now the executable isn't enough! :')
- prevailing approach to dynamic linking uses machine code as the “lowest common denominator” aka linking code only at the machine level

Loader

General

- Input = executable code
- output = running programs
- executable files are stored on disk
- loader's job is essentially to load the file onto memory and start running it (aka loader = OS)

How it works

- Reads executable file's header to determine size of text and data segments
- Creates new address space for program large enough to hold text and data segments, along with a stack segment
- Copies instructions + data from executable file into the new address space
- Copies arguments passed to the program onto the stack
- Initializes machine registers
 - Most registers cleared, but stack pointer assigned address of 1st free stack location
- Jumps to start-up routine that copies program's arguments from stack to registers & sets the PC
 - If main routine returns, start-up routine terminates program with exit system call

SDS/Boolean Algebra/FSM

Hold Time

$$t_{\text{clk-to-q}} + t_{\text{logic, shortest}} > t_{\text{hold}}$$

- look for the shortest combinatorial path
- the longest that the path can be stable